

Application Performance Tuning

M. D. Jones, Ph.D.

Center for Computational Research
University at Buffalo
State University of New York

High Performance Computing I, 2013

Performance Foundations

Three pillars of performance optimization:

Algorithmic - choose the most effective algorithm that you can for the problem of interest

Serial Efficiency - optimize the code to run efficiently in a non-parallel environment

Parallel Efficiency - effectively use multiple processors to achieve a reduction in execution time, or equivalently, to solve a proportionately larger problem

Algorithmic Efficiency

Choose the best algorithm *before* you start coding (recall that good planning is an essential part of writing good software):

- Running on large number of processors? Choose an algorithm that **scales** well with increasing processor count
- Running a large system (mesh points, particle count, etc.)? Choose an algorithm that **scales** well with system size
- If you are going to run on a massively parallel machine, plan from the beginning on how you intend to decompose the problem (it may save you a lot of time later)

Serial Efficiency

Getting efficient code in parallel is made much more difficult if you have not optimized the sequential code, and in fact can lead to a misleading picture of parallel performance. Recall that our definition of parallel speedup,

$$S(N_p) = \frac{\tau_S}{\tau_p(N_p)},$$

involves the time, τ_S , for an **optimal** sequential implementation (**not** just $\tau_p(1)$!)

Establish a Performance Baseline

Steps to establishing a baseline for your own performance expectations:

- Choose a representative problem (or better still a suite of problems) that can be run under identical circumstances on a multitude of platforms/compilers

Requires portable code!

- How fast is “fast”? You can utilize hardware performance counters to measure actual code performance
- Profile, profile, and then profile some more ... to find bottlenecks and spend **your** time more effectively in optimizing code

Parallel Performance Trap

Pitfalls when measuring the performance of parallel codes:

- For many, *speedup* or *linear scalability* is the ultimate goal.
- This goal is incomplete - a terribly inefficient code can scale well, but actually deliver poor **efficiency**.
- For example, consider a simple Monte Carlo based code that uses the most rudimentary uniform sampling (i.e. no importance sampling) - this can be made to scale perfectly in parallel, but the algorithmic efficiency (measured perhaps by the delivered variance per cpu-hour) is quite low.

time Command

Note that this is not the `time` built-in function in many shells (`bash` and `tcsh` included), but instead the one located in `/usr/bin`. This command is quite useful for getting an overall picture of code performance. The default output format:

```
%User %System %Elapsed %PCPU (%Xtext+%Ddata %Mmax)k
%Iinputs+%Ooutputs (%Fmajor+%Rminor)pagefaults %Wswaps
```

and using the `-p` option:

```
real %e
user %U
sys %S
```

time Example

```
[rush:~/d_laplace]$ /usr/bin/time ./laplace_s
:
  Max value in sol:   0.999992327961218
  Min value in sol:  -8.742278000372475E-008
75.82user 0.00system 1:17.72elapsed 97%CPU (0avgtext+0avgdata 0maxresident)k
0inputs+0outputs (0major+913minor)pagefaults 0swaps
[bono:~/d_laplace]$ /usr/bin/time -p ./laplace_s
:
  Max value in sol:   0.999992327961218
  Min value in sol:  -8.742278000372475E-008
real 75.73
user 74.68
sys 0.00
```


time MPI Example

```
[rush:~/d_laplace]$ /usr/bin/time mpiexec -np 2 ./laplace_mpi
:
  Max value in sol:   0.999992327961218
  Min value in sol:  -8.742278000372475E-008
Writing logfile....
Finished writing logfile.
28.43user 1.54system 0:31.95elapsed 93%CPU (0avgtext+0avgdata 0maxresident)k
0inputs+0outputs (0major+14920minor)pagefaults 0swaps
```

You will see the result of timing the `mpiexec` shell script, not the MPI code). Of course, having external timing is nice, but thankfully MPI gives up much better timing and profiling tools to use.

Code Section Timing (Calipers)

Timing sections of code requires a bit more work on the part of the programmer, but there are *reasonably* portable means of doing so:

Routine	Type	Resolution
times	user/sys	ms
gettimeofday	wall	μ s
clock_gettime	wall	ns
system_clock (f90)	wall	system-dependent
cpu_time (f95)	cpu	compiler-dependent
MPI_Wtime*	wall	system-dependent
OMP_GET_WTIME*	wall	system-dependent

Generally I prefer the MPI and OpenMP timing calls whenever I can use them (*the MPI and OpenMP specifications call for their intrinsic timers to be high precision).

More information on code section timers (and code for doing so):

- LLNL Performance Tools:

https://computing.llnl.gov/tutorials/performance_tools/#gettimeofday

- Stopwatch (nice F90 module, but you need to supply a low-level function for accessing a timer):

<http://math.nist.gov/StopWatch>

GNU Tools: gprof

Tool that we used briefly before:

- Generic GNU profiler
- Requires recompiling code with **-pg** option
- Running subsequent instrumented code produces **gmon.out** to be read by **gprof**
- Use the environment variable **GMON_OUT_PREFIX** to specify a new **gmon.out** prefix to which the process ID will be appended (especially useful for parallel runs) - this is a largely undocumented feature ...
- Line-level profiling is possible, as we will see in the following example

gprof Shortcomings

Shortcomings of **gprof** (which apply also to any statistical profiling tool):

- Need to recompile to instrument the code
- Instrumentation can affect the statistics in the profile
- Overhead can significantly increase the running time
- Compiler optimization can be affected by instrumentation

Types of gprof Profiles

gprof profiles come in three types:

- 1 **Flat Profile:** shows how much time your program spent in each function, and how many times that function was called
- 2 **Call Graph:** for each function, which functions called it, which other functions it called, and how many times. There is also an estimate of how much time was spent in the subroutines of each function
- 3 **Basic-block:** Requires compilation with the `-a` flag (supported only by GNU?) - enables gprof to construct an annotated source code listing showing how many times each line of code was executed

gprof example

```

gfortran -I. -O3 -ffast-math -g -pg -o rp rp_read.o initial.o en_gde.o \
        adwfns.o rpqmc.o evol.o dbxgde.o dbxtri.o gen_etg.o gen_rtg.o \
        gdewfn.o lib/*.o
[jonesm@rush ~/d_bench]$ qsub -qdebug -lnodes=1:ppn=2,walltime=00:30:00 -I
[jonesm@rush ~/d_bench]$ ./rp
Enter runid:(<=9chars)
short
...
...skip copious amount of standard output ...
...
[jonesm@rush ~/d_bench]$ gprof rp gmon.out > & out.gprof
[jonesm@rush ~/d_bench]$ less out.gprof
Flat profile:

Each sample counts as 0.01 seconds.
 %   cumulative   self           self      total
time  seconds    seconds   calls   s/call   s/call   name
89.74    123.23    123.23    204008     0.00     0.00   triwfns_
 6.96     132.79     9.56         1     9.56   137.22   MAIN__
 1.18     134.41     1.62    200004     0.00     0.00   en_gde__
 1.05     135.86     1.44    204002     0.00     0.00   evol_
 0.71     136.83     0.97  14790551     0.00     0.00   ranf_
 0.27     137.20     0.37    204008     0.00     0.00   gdewfn_
...

```

```
[jonesm@rush ~/d_bench]$ gprof --line rp gmon.out > & out.gprof.line
[jonesm@rush ~/d_bench]$ less out.gprof.line
Flat profile:
```

Each sample counts as 0.01 seconds.

%	cumulative	self		self	total	
time	seconds	seconds	calls	ns/call	ns/call	name
17.45	23.96	23.96				triwfns_ (adwfns.f:129 @ 403c94)
14.46	43.82	19.86				triwfns_ (adwfns.f:130 @ 403ce6)
12.87	61.50	17.68				triwfns_ (adwfns.f:171 @ 404755)
12.31	78.41	16.91				triwfns_ (adwfns.f:172 @ 4047e2)
0.67	79.33	0.92				triwfns_ (adwfns.f:130 @ 403cd6)
0.59	80.14	0.82				MAIN__ (cc4WTuQH.f:308 @ 4070c6)
0.51	80.84	0.70				MAIN__ (cc4WTuQH.f:304 @ 4072da)

More gprof Information

More **gprof** documentation:

- **gprof** GNU Manual:

<http://sourceware.org/binutils/docs/gprof/>

- **gprof** man page:

[man gprof](#)


PGI Tools: pgprof

- PGI tools also have profiling capabilities (c.f. **man pgf95**)
- Graphical profiler, **pgprof**

pgprof example

```
pgf90 -tp p7-64 -fastsse -g77libs -Mprof=lines -o rp rp_read.o initial.o en_gde.o \  
      adwfns.o rpqmc.o evol.o dbxgde.o dbxtri.o gen_etg.o gen_rtg.o \  
      gdewfn.o lib/*.o  
[jonesm@rush ~/d_bench]$ ./rp  
Enter runid:(<=9chars)  
short  
...  
...skip copious amount of standard output ...  
...  
[jonesm@rush ~/d_bench]$ pgprof -exe ./rp pgprof.out
```

pgprof example [screenshot]



- N.B. you can also use the **-text** option to **pgprof** to make it behave more like **gprof**
- See the *PGI Tools Guide* for more information (should be a PDF copy in \$PGI/doc)

mpiP: Statistical MPI Profiling

- <http://mpip.sourceforge.net>
- Not a tracing tool, but a lightweight interface to accumulate statistics using the MPI profiling interface
- Quite useful in conjunction with a tracefile analysis (e.g. using jumpshot)
- Installed on CCR systems - see "**module avail**" for **mpiP** availability and location

mpiP Compilation

To use mpiP you need to:

- Add a `-g` flag to add symbols (this will allow mpiP to access the source code symbols and line numbers)
- Link in the necessary mpiP profiling library and the binary utility libraries for actually decoding symbols (*There is a trick that you can use most of the time to avoid having to link with mpiP, though*).

Compilation examples (from U2) follow ...

mpiP Runtime Flags

You can set various mpiP runtime flags (e.g. `export MPIP="-t 10.0 -k 2"`):

Option	Description	Default
-c	Generate concise version of report, omitting callsite process-specific detail.	
-e	Print report data using floating-point format	
-f dir	Record output file in directory <dir>	.
-g	Enable mpiP debug mode	disabled
-k n	Sets callsite stack traceback depth to <n>	1
-n	Do not truncate full pathname of filename in callsites	
-o	Disable profiling at initialization. Application must enable profiling with <code>MPI_Pcontrol()</code>	
-s n	Set hash table size to <n>	256
-t x	Set print threshold for report, where <x> is the MPI percentage of time for each callsite	0.0
-v	Generates both concise and verbose report output	

mpiP Example Output

From an older version of mpiP, but still almost entirely the same - this one links directly with mpiP; first compile:

```
mpicc -g -o atlas aij2_basis.o analyze.o atlas.o barrier.o byteflip.o \  
chordgn2.o cstrings.o io2.o map.o mutils.o numrec.o paramods.o proj.o \  
projAtlas.o sym2.o util.o -lm -L/Projects/CCR/jonesm/mpiP-2.8.2/gnu/ch_gm/lib \  
-lmpiP -lbfd -liberty -lm
```

then run (in this case using 16 processors) and examine the output file:

```
[jonesm@joplin d_derenzo]$ ls *.mpiP  
atlas.gcc3-mpi-papi-mpiP.16.20578.1.mpiP  
[jonesm@joplin d_derenzo]$ less atlas.gcc3-mpi-papi-mpiP.16.20578.1.mpiP  
:  
:
```

```
@ mpiP
@ Command : ./atlas.gcc3-mpi-papi-mpiP study.ini 0
@ Version : 2.8.2
@ MPIP Build date : Jun 29 2005, 14:53:41
@ Start time : 2005 06 29 15:18:52
@ Stop time : 2005 06 29 15:28:34
@ Timer Used : gettimeofday
@ MPIP env var : [null]
@ Collector Rank : 0
@ Collector PID : 20578
@ Final Output Dir : .
@ MPI Task Assignment : 0 bb18n17.ccr.buffalo.edu
@ MPI Task Assignment : 1 bb18n17.ccr.buffalo.edu
@ MPI Task Assignment : 2 bb18n16.ccr.buffalo.edu
@ MPI Task Assignment : 3 bb18n16.ccr.buffalo.edu
@ MPI Task Assignment : 4 bb18n15.ccr.buffalo.edu
@ MPI Task Assignment : 5 bb18n15.ccr.buffalo.edu
@ MPI Task Assignment : 6 bb18n14.ccr.buffalo.edu
@ MPI Task Assignment : 7 bb18n14.ccr.buffalo.edu
@ MPI Task Assignment : 8 bb18n13.ccr.buffalo.edu
@ MPI Task Assignment : 9 bb18n13.ccr.buffalo.edu
@ MPI Task Assignment : 10 bb18n12.ccr.buffalo.edu
@ MPI Task Assignment : 11 bb18n12.ccr.buffalo.edu
@ MPI Task Assignment : 12 bb18n11.ccr.buffalo.edu
@ MPI Task Assignment : 13 bb18n11.ccr.buffalo.edu
@ MPI Task Assignment : 14 bb18n10.ccr.buffalo.edu
@ MPI Task Assignment : 15 bb18n10.ccr.buffalo.edu
```

```

-----
@--- MPI Time (seconds) -----
-----

```

Task	AppTime	MPITime	MPI%
0	582	44.7	7.69
1	579	41.9	7.24
2	579	40.7	7.03
3	579	36.9	6.37
4	579	22.3	3.84
5	579	16.6	2.87
6	579	32	5.53
7	579	35.9	6.20
8	579	28.6	4.93
9	579	25.9	4.48
10	579	39.2	6.76
11	579	33.8	5.84
12	579	35.3	6.10
13	579	41	7.07
14	579	29.9	5.16
15	579	41.4	7.16
*	9.27e+03	546	5.89

```
-----
@--- Callsites: 13 -----
-----
```

ID	Lev	File/Address	Line	Parent_Funct	MPI_Call
1	0	util.c	833	gsync	Barrier
2	0	atlas.c	1531	readProjData	Allreduce
3	0	projAtlas.c	745	backProjAtlas	Allreduce
4	0	atlas.c	1545	readProjData	Allreduce
5	0	atlas.c	1525	readProjData	Allreduce
6	0	atlas.c	1541	readProjData	Allreduce
7	0	atlas.c	1589	readProjData	Allreduce
8	0	atlas.c	1519	readProjData	Allreduce
9	0	util.c	789	mygcast	Bcast
10	0	projAtlas.c	1100	computeLoglikeAtlas	Allreduce
11	0	atlas.c	1514	readProjData	Allreduce
12	0	atlas.c	1537	readProjData	Allreduce
13	0	projAtlas.c	425	fwdBackProjAtlas2	Allreduce

```

-----
@--- Aggregate Time (top twenty, descending, milliseconds) -----
-----
Call           Site      Time      App%      MPI%      COV
Allreduce      13      3.09e+05  3.33      56.50     0.46
Barrier        1       2.13e+05  2.30      38.97     0.35
Bcast          9       1.69e+04  0.18      3.10      0.37
Allreduce      3       7.78e+03  0.08      1.42      0.11
Allreduce      10      62.7      0.00      0.01      0.20
Allreduce      11      2.42      0.00      0.00      0.09
Allreduce      7       2.17      0.00      0.00      0.26
Allreduce      12      1.15      0.00      0.00      0.20
Allreduce      6       1.14      0.00      0.00      0.19
Allreduce      5       1.13      0.00      0.00      0.15
Allreduce      8       1.12      0.00      0.00      0.18
Allreduce      2       1.1       0.00      0.00      0.13
Allreduce      4       1.1       0.00      0.00      0.12

```

```

-----
@--- Aggregate Sent Message Size (top twenty, descending, bytes) -----
-----

```

Call	Site	Count	Total	Avrg	Sent%
Allreduce	13	65536	2.28e+09	3.48e+04	83.69
Allreduce	3	8192	2.85e+08	3.48e+04	10.46
Bcast	9	490784	1.59e+08	325	5.85
Allreduce	11	16	2.07e+04	1.3e+03	0.00
Allreduce	10	512	4.1e+03	8	0.00
Allreduce	7	16	256	16	0.00
Allreduce	2	16	64	4	0.00
Allreduce	6	16	64	4	0.00
Allreduce	5	16	64	4	0.00
Allreduce	4	16	64	4	0.00
Allreduce	8	16	64	4	0.00
Allreduce	12	16	64	4	0.00
...					
...					

Using mpiP at Runtime

Now let's examine an example using mpiP at runtime. This example solves a simple Laplace equation with Dirichlet boundary conditions using finite differences.

```
1  #!/bin/bash
2  #SBATCH --nodes=2
3  #SBATCH --ntasks-per-node=8
4  #SBATCH --constraint=CPU-L5520
5  #SBATCH --partition=debug
6  #SBATCH --time=00:10:00
7  #SBATCH --mail-type=END
8  #SBATCH --mail-user=jonesm@buffalo.edu
9  #SBATCH --output=slurmMPIP.out
10 #SBATCH --job-name=mpip-test
11 module load intel
12 module load intel-mpi
13 module load mpip
14 module list
15 export I_MPI_DEBUG=4
16 # Use LD_PRELOAD trick to load mpiP wrappers at runtime
17 export LD_PRELOAD=$MPIDIR/lib/libmpiP.so
18 export I_MPI_PMI_LIBRARY=/usr/lib64/libpmi.so
19 srun ./laplace_mpi<<EOF
20 2000
21 EOF
```


... and then run it and examine the resulting mpiP output file:

```
[rush:~/d_laplace/d_mpip]$ ls -l laplace_mpi.*
[rush:~/d_laplace/d_mpip]$ ls -l laplace_mpi.*
-rw-r--r-- 1 jonesm ccrstaff 17920 Dec 5 2012 laplace_mpi.16.11597.1.mpiP
-rw-r--r-- 1 jonesm ccrstaff 17698 Oct 7 15:45 laplace_mpi.16.31074.1.mpiP
-rw-r--r-- 1 jonesm ccrstaff 16032024 Oct 7 15:45 laplace_mpi.dat
lrwxrwxrwx 1 jonesm ccrstaff 18 Dec 21 2010 laplace_mpi.f90 -> ../laplace_mpi.f90
```

and again we will break it down by section:

```
1 @ mpiP
2 @ Command : /ifs/user/jonesm/d_laplace/d_mpiP/./laplace_mpi
3 @ Version : 3.3.0
4 @ MPIP Build date : Oct 14 2011, 16:16:34
5 @ Start time : 2013 10 07 15:43:14
6 @ Stop time : 2013 10 07 15:45:11
7 @ Timer Used : PMPI_Wtime
8 @ MPIP env var : [null]
9 @ Collector Rank : 0
10 @ Collector PID : 31074
11 @ Final Output Dir : .
12 @ Report generation : Single collector task
13 @ MPI Task Assignment : 0 d16n02
14 @ MPI Task Assignment : 1 d16n02
15 @ MPI Task Assignment : 2 d16n02
16 @ MPI Task Assignment : 3 d16n02
17 @ MPI Task Assignment : 4 d16n02
18 @ MPI Task Assignment : 5 d16n02
19 @ MPI Task Assignment : 6 d16n02
20 @ MPI Task Assignment : 7 d16n02
21 @ MPI Task Assignment : 8 d16n03
22 @ MPI Task Assignment : 9 d16n03
23 @ MPI Task Assignment : 10 d16n03
24 @ MPI Task Assignment : 11 d16n03
25 @ MPI Task Assignment : 12 d16n03
26 @ MPI Task Assignment : 13 d16n03
27 @ MPI Task Assignment : 14 d16n03
28 @ MPI Task Assignment : 15 d16n03
```

```

29 -----
30 @--- MPI Time (seconds) ---
31 -----
32 Task      AppTime      MPITime      MPI%
33   0         117         17.2        14.70
34   1         117         17.6        15.08
35   2         117         19.9        17.03
36   3         117         17.9        15.33
37   4         117         17.7        15.14
38   5         117         14.6        12.53
39   6         117         13.6        11.64
40   7         117         13.4        11.49
41   8         117         12.2        10.48
42   9         117         16.7        14.28
43  10         117          17         14.52
44  11         117         17.2        14.70
45  12         117         19.1        16.33
46  13         117         17.5        14.97
47  14         117         18.3        15.68
48  15         117         17.5        14.96
49  *      1.87e+03         267        14.30

```

```

50 -----
51 @--- Callsites: 6 -----
52 -----
53 ID Lev File/Address      Line Parent_Funct      MPI_Call
54 1  0 laplace_mpi.f90      118 MAIN__           Allreduce
55 2  0 laplace_mpi.f90      143 MAIN__           Recv
56 3  0 laplace_mpi.f90      48  __paramod_MOD_xchange Sendrecv
57 4  0 laplace_mpi.f90      80 MAIN__           Bcast
58 5  0 laplace_mpi.f90      46  __paramod_MOD_xchange Sendrecv
59 6  0 laplace_mpi.f90      138 MAIN__           Send
60 -----
61 @--- Aggregate Time (top twenty, descending, milliseconds) -----
62 -----
63 Call          Site      Time      App%      MPI%      COV
64 Allreduce      1      2.45e+05  13.09     91.53     0.14
65 Sendrecv       5      1.12e+04   0.60      4.20      0.17
66 Sendrecv       3      1.11e+04   0.59      4.14      0.20
67 Send           6          306      0.02      0.11      0.55
68 Bcast          4          16.7     0.00      0.01      0.25
69 Recv           2           14      0.00      0.01      0.00

```

```

70 -----
71 @--- Callsite Time statistics (all, milliseconds): 80 -----
72 -----
73 Name                Site Rank  Count      Max      Mean      Min      App%      MPI%
74 Allreduce            1      0  23845    32.2    0.667    0.037    13.57    92.36
75 Allreduce            1      1  23845    29.1    0.685    0.037    13.98    92.72
76 Allreduce            1      2  23845    33.9    0.779    0.037    15.92    93.45
77 Allreduce            1      3  23845    29.2    0.696     0.04    14.22    92.77
78 Allreduce            1      4  23845    32.3    0.687    0.036    14.03    92.65
79 Allreduce            1      5  23845    29.2     0.56    0.037    11.44    91.29
80 Allreduce            1      6  23845    28.5    0.511    0.035    10.43    89.60
81 Allreduce            1      7  23845    25.7    0.493    0.036    10.06    87.62
82 Allreduce            1      8  23845    26.1    0.444    0.038     9.07    86.55
83 Allreduce            1      9  23845    28.2    0.637    0.034    13.00    91.04
84 Allreduce            1     10  23845    33.6    0.649    0.036    13.26    91.30
85 Allreduce            1     11  23845    28.9    0.657    0.034    13.42    91.29
86 Allreduce            1     12  23845    32.3    0.737    0.039    15.06    92.23
87 Allreduce            1     13  23845     31     0.67    0.035    13.69    91.46
88 Allreduce            1     14  23845    33.7    0.704    0.038    14.38    91.73
89 Allreduce            1     15  23845    33.9    0.683    0.036    13.94    93.18
90 Allreduce            1      * 381520    33.9    0.641    0.034    13.09    91.53

```

```

91 -----
92 @--- Callsite Message Sent statistics (all, sent bytes) ---
93 -----
94 Name           Site Rank   Count      Max      Mean      Min      Sum
95 Allreduce      1      0   23845      8        8        8 1.908e+05
96 Allreduce      1      1   23845      8        8        8 1.908e+05
97 Allreduce      1      2   23845      8        8        8 1.908e+05
98 Allreduce      1      3   23845      8        8        8 1.908e+05
99 Allreduce      1      4   23845      8        8        8 1.908e+05
100 Allreduce     1      5   23845      8        8        8 1.908e+05
101 Allreduce     1      6   23845      8        8        8 1.908e+05
102 Allreduce     1      7   23845      8        8        8 1.908e+05
103 Allreduce     1      8   23845      8        8        8 1.908e+05
104 Allreduce     1      9   23845      8        8        8 1.908e+05
105 Allreduce     1     10   23845      8        8        8 1.908e+05
106 Allreduce     1     11   23845      8        8        8 1.908e+05
107 Allreduce     1     12   23845      8        8        8 1.908e+05
108 Allreduce     1     13   23845      8        8        8 1.908e+05
109 Allreduce     1     14   23845      8        8        8 1.908e+05
110 Allreduce     1     15   23845      8        8        8 1.908e+05
111 Allreduce     1      * 381520      8        8        8 3.052e+06

```

Intel Trace Analyzer/Collector (ITAC)

A commercial product for performing MPI trace analysis that has enjoyed a long history is **Vampir/Vampirtrace**, originally developed and sold by Pallas GmbH. Now owned by Intel and available as the Intel Trace Analyzer and Collector. We have a license on U2 if someone wants to give it a try.

Note that **Vampir/Vampirtrace** has since been reborn as an entirely new product.

ITAC Example

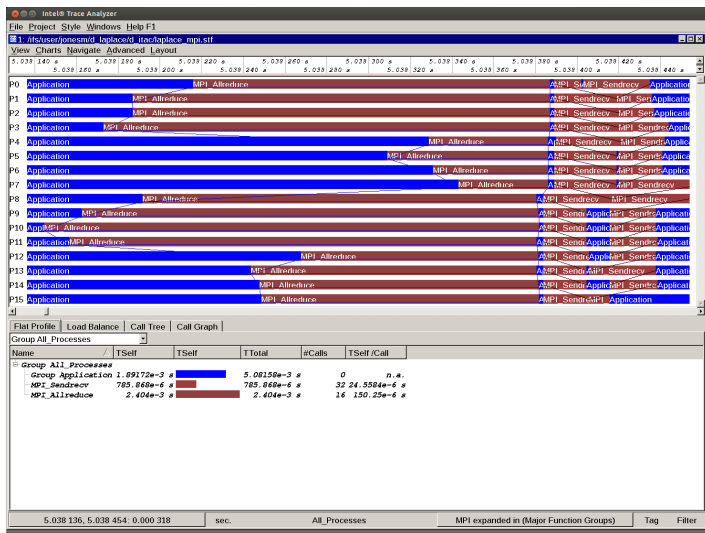
Note that you do not have to recompile your application to use ITAC (unless you are building it statically), you can just build it usual, using Intel MPI:

```
[rush:~/d_laplace/d_itac]$ module load intel-mpi
[rush:~/d_laplace/d_itac]$ make laplace_mpi
mpiifort -ipo -O3 -Vaxlib -g -c laplace_mpi.f90
mpiifort -ipo -O3 -Vaxlib -g -o laplace_mpi laplace_mpi.o
ipo: remark #11001: performing single-file optimizations
ipo: remark #11005: generating object file /tmp/ipo_ifortjOzYAn.o
[bono:~/d_laplace/d_itac]$
```

You can turn trace collection on at run-time ...


```
1  #!/bin/bash
2  #SBATCH --nodes=2
3  #SBATCH --ntasks-per-node=8
4  #SBATCH --constraint=CPU-L5520
5  #SBATCH --partition=debug
6  #SBATCH --time=00:10:00
7  #SBATCH --mail-type=END
8  #SBATCH --mail-user=jonesm@buffalo.edu
9  #SBATCH --output=slurmITAC.out
10 #SBATCH --job-name=mpip-test
11 module load intel
12 module load intel-mpi
13 . /util/intel/itac/8.1.0.024/bin/itacvars.sh
14 module list
15 which mpiexec
16 # this is one of those times when srun is not going to work - need to
17 # use Intel MPI's task launching to generate ITAC profile
18 MYHOSTFILE=tmp.hosts
19 srun -l hostname -s | sort -n | awk '{print $2}' > $MYHOSTFILE
20 NNODES=`cat $MYHOSTFILE | uniq | wc -l`
21 NPROCS=`cat $MYHOSTFILE | wc -l`
22 export I_MPI_DEBUG=4
23 mpdboot -n $NNODES -f "$MYHOSTFILE" -v
24 mpdtrace
25 mpiexec -trace -np $NPROCS ./laplace_mpi <<EOF
26 2000
27 EOF
28 mpdallexit
29 [ -e "$MYHOSTFILE" ] && \rm "$MYHOSTFILE"
```

Running the preceding batch job on U2 produces a bunch (many!) of profiling output files, the most important of which can be is the name of your binary with a **.stf** suffix, in this case *laplace_mpi.stf*, which we feed to the Intel Trace Analyzer using the **traceanalyzer** command ... and we should see a profile that looks very much like what you can see using **jumpshot** (MPICH2's trace file profiler).



More ITAC Documentation

Some helpful pointers to more ITAC documentation:

```
1 [rush:~/d_laplace/d_itac]$ which traceanalyzer
2 /util/intel/itac/8.1.0.024/bin/traceanalyzer
3 [rush:~/d_laplace/d_itac]$ ls -l /util/intel/itac/8.1.0.024/doc
4 FAQ.pdf
5 Getting_Started.html
6 INSTALL.html
7 ITA_Reference_Guide
8 ITA_Reference_Guide.htm
9 ITA_Reference_Guide.pdf
10 ITC_Reference_Guide
11 ITC_Reference_Guide.htm
12 ITC_Reference_Guide.pdf
13 Release_Notes_Addendum_for_MIC_Architecture.txt
14 Release_Notes.txt
```

Generally a good idea to refer to the documentation for the same version that you are using (you can check with **"module show intel-mpi"**).

Introduction

- **P**erformance **A**pplication **P**rogramming Interface
- Implement a portable(!) and efficient API to access existing **hardware** performance counters
- Ease the optimization of code by providing base infrastructure for cross-platform tools

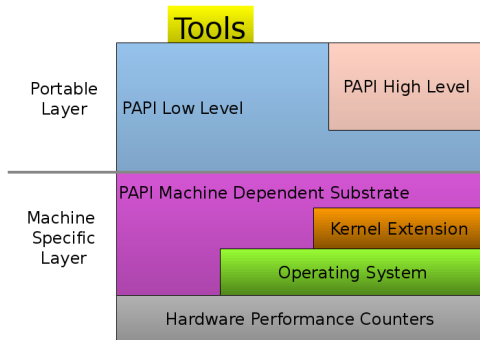
Pre-PAPI

Before PAPI came along, there were hardware performance counters, of course - but access to them was limited to proprietary tools and APIs. Some examples were SGI's **perfex** and Cray's **hpm**.

Now, as long as PAPI has been ported to a particular hardware substrate, the end-programmer (or tool developer) can just use the PAPI interface.

PAPI Schematic

Best summarized by the following schematic picture:



Behind the PAPI Curtain

- Linux - x86/x86_64 uses the **perfctr** kernel patches by Mikael Petterssen:

<http://user.it.uu.se/~mikpe/linux/perfctr/2.6>

- Headed for inclusion in mainstream Linux kernel (was a custom patch applied to CCR systems prior to Linux kernel 2.6.32)
 - low overhead
- IA64 - uses **PFM**, developed by HP and included in the linux kernel (for x86_64):
 - Full use of available IA64 monitoring capabilities
 - Quite a bit slower than **perfctr**, at least according to the **PAPI** developers
 - <http://www.hpl.hp.com/research/linux/perfmon>
 - libpfm lives on using perf events, but perfmon apparently ceased development for Linux as of kernel 2.6.30 or so
- "Perf Events" added to Linux kernel in 2.6.31, replacing both of the above, c.f.:

<http://web.eecs.utk.edu/~vweaver1/projects/perf-events/>

Block diagram for Nehalem architecture, showing a single socket (repeat on QPI for dual sockets):



“Westmere” Xeons

Characteristics of the ‘Westmere’ E5645 Xeons that form (part of) CCR’s **U2** cluster:

Clock Cycle	2.4 GHz
TPP	9.6 GFlop/s (per core)
Pipeline	14 stages
L2 Cache Size	256 kByte
L3 Cache Size	12 MByte
CPU-Memory Bandwidth	32 GByte/s (nonuniform!)

“Westmere” Xeon Memory Hierarchy Penalties

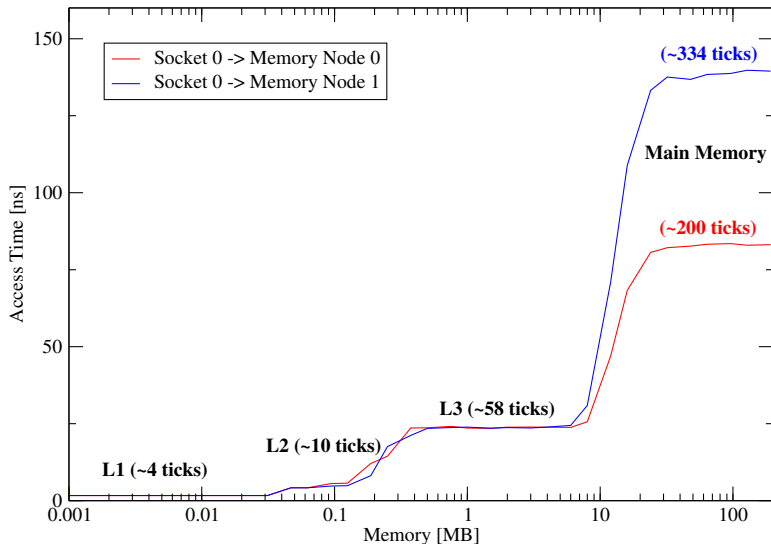
Consider the penalties in lost computation for **not** reusing data in the various caches (using U2's Intel *Westmere* Xeon E5645 processors):

Memory	Miss Penalty (cycles)
L1 Cache	4
L2 Cache	10
L3 Cache	58
Main	200 (on socket) 334 (off socket)

as determined by the **lmbench** benchmark¹.

¹<http://www.bitmover.com/lmbench>

Xeon E5645 Memory Latencies



Available PAPI Performance Data

- Cycle count
- Instruction count (including Integer, Floating point, load/store)
- Branches (including Taken/not taken, mispredictions)
- Pipeline stalls (due to memory, resource conflicts)
- Cache (misses for different levels, invalidation)
- TLB (Translation Lookaside Buffer) misses, invalidation, etc.

High-level PAPI

- Intended for coarse-grained measurements
- Requires little (or no) setup code
- Allows only PAPI preset
- Allows only aggregate counting (no statistical profiling)

Low-level PAPI

- More efficient (and functional) than high-level
- About 60 functions
- Thread-safe
- Supports presets and native events

Preset vs. Native Events

preset or pre-defined events, are those which have been considered useful by the PAPI community and developers:

<http://icl.cs.utk.edu/projects/papi/presets.html>

native events are those countable by the CPU's hardware. These events are highly platform specific, and you would need to consult the processor architecture manuals for the relevant native event lists

Low-level PAPI Functions

- Hardware counter multiplexing (time sharing hardware counters to allow more events to be monitored than can be conventionally supported)
- Processor information
- Address space information
- Memory information (static and dynamic)
- Timing functions
- Hardware event inquiry
- ... and many more

More PAPI Information

For more on PAPI, including source code, documentation, presentations, and links to third-party tools that utilize PAPI, see

<http://icl.cs.utk.edu/projects/papi>

How to Access PAPI at CCR

Consider a simple example code to measure Flop/s using the high-level PAPI API:

```
1  #include <stdio.h>
2  #include <stdlib.h>
3  #include "papi.h"
4
5
6  int main()
7  {
8      float real_time, proc_time, mflops;
9      long_long flpops;
10     float ival_time, ival_time, imflops;
11     long_long iflpops;
12     int retval;
```

```
13  if((retval=PAPI_flops(&iREAL_time,&iPROC_time,&iFLPOPS,&iMFLPOPS)) < PAPI_OK)
14  {
15      printf("Could not initialise PAPI_flops \n");
16      printf("Your platform may not support floating point operation event.\n");
17      printf("retval: %d\n", retval);
18      exit(1);
19  }
20
21  your_slow_code();
22
23  if((retval=PAPI_flops( &real_time, &proc_time, &flpops, &mflops))<PAPI_OK)
24  {
25      printf("retval: %d\n", retval);
26      exit(1);
27  }
28
29  printf("Real_time: %f Proc_time: %f Total flpops: %lld MFLOPS: %f\n",
30         real_time, proc_time, flpops, mflops);
31  exit(0);
32 }
33 int your_slow_code()
34 {
35     int i;
36     double tmp=1.1;
37
38     for(i=1; i<2000; i++)
39     {
40         tmp=(tmp+100)/i;
41     }
42     return 0;
43 }
```

How to Access PAPI at CCR

On U2, you access the **papi** module and compile accordingly:

```
[rush:/ifs/user/jonesm/d_papi]$ module list
Currently Loaded Modulefiles:
  1) null                2) modules              3) use.own              4) intel-mpi/4.1.1
  5) papi/v5.1.1         6) intel/13.1
[rush:~/d_papi]$ icc -I$PAPI/include -o PAPI_flops PAPI_flops.c -L$PAPI/lib -lpapi
[rush:~/d_papi]$ ./PAPI_flops
Real_time: 0.000017 Proc_time: 0.000002 Total flpops: 16 MFLOPS: 6.636506
```

... and on Lennon (SGI Altix):

```
[jonesm@lennon ~/d_papi]$ module load papi
'papi/v3.2.1' load complete.
[jonesm@lennon ~/d_papi]$ echo $PAPI
/util/perftools/papi-3.2.1
[jonesm@lennon ~/d_papi]$ gcc -I$PAPI/include -o PAPI_flops PAPI_flops.c \
-L$PAPI/lib -lpapi
[jonesm@lennon ~/d_papi]$ ldd PAPI_flops
    libpapi.so => /util/perftools/papi-3.2.1/lib/libpapi.so
                                   (0x2000000000040000)
    libc.so.6.1 => /lib/tls/libc.so.6.1 (0x20000000000ac000)
    libpfm.so.2 => /usr/lib/libpfm.so.2 (0x2000000000318000)
    /lib/ld-linux-ia64.so.2 => /lib/ld-linux-ia64.so.2 (0x2000000000000000)
[jonesm@lennon ~/d_papi]$ ./PAPI_flops
Real_time: 0.000143 Proc_time: 0.000132 Total flpops: 48056 MFLOPS: 363.964111
```

N.B., The Altix is long dead, but this is a good example of the cross-platform portability of PAPI accessign the hardware performance counters.

papi_avail Command

You can use **papi_avail** to check event availability (different CPUs support various events):

```
1  #!/bin/bash
2  #SBATCH --nodes=1
3  #SBATCH --ntasks-per-node=8
4  #SBATCH --constraint=CPU-L5520
5  #SBATCH --partition=debug
6  #SBATCH --time=00:10:00
7  #SBATCH --mail-type=END
8  #SBATCH --mail-user=jonesm@buffalo.edu
9  #SBATCH --output=slurmQ.out
10 #SBATCH --job-name=papi-test
11 #
12 module load papi
13 module list
14 export | grep SLURM
15 papi_avail
```

Available events and hardware information.

```
-----  
PAPI Version           : 5.1.1.0  
Vendor string and code : GenuineIntel (1)  
Model string and code  : Intel(R) Xeon(R) CPU           L5520  @ 2.27GHz (26)  
CPU Revision           : 5.000000  
CPUID Info             : Family: 6  Model: 26  Stepping: 5  
CPU Max Megahertz      : 2266  
CPU Min Megahertz      : 2266  
Hdw Threads per core   : 1  
Cores per Socket       : 4  
NUMA Nodes             : 2  
CPUs per Node          : 4  
Total CPUs             : 8  
Running in a VM        : no  
Number Hardware Counters : 7  
Max Multiplex Counters  : 64  
-----
```


Name	Code	Avail	Deriv	Description (Note)
PAPI_L1_DCM	0x80000000	Yes	No	Level 1 data cache misses
PAPI_L1_ICM	0x80000001	Yes	No	Level 1 instruction cache misses
PAPI_L2_DCM	0x80000002	Yes	Yes	Level 2 data cache misses
PAPI_L2_ICM	0x80000003	Yes	No	Level 2 instruction cache misses
PAPI_L3_DCM	0x80000004	No	No	Level 3 data cache misses
PAPI_L3_ICM	0x80000005	No	No	Level 3 instruction cache misses
PAPI_L1_TCM	0x80000006	Yes	Yes	Level 1 cache misses
PAPI_L2_TCM	0x80000007	Yes	No	Level 2 cache misses
PAPI_L3_TCM	0x80000008	Yes	No	Level 3 cache misses
PAPI_CA_SNP	0x80000009	No	No	Requests for a snoop
PAPI_CA_SHR	0x8000000a	No	No	Requests for exclusive access to shared cache line
PAPI_CA_CLN	0x8000000b	No	No	Requests for exclusive access to clean cache line
PAPI_CA_INV	0x8000000c	No	No	Requests for cache line invalidation
PAPI_CA_ITV	0x8000000d	No	No	Requests for cache line intervention
PAPI_L3_LDM	0x8000000e	Yes	No	Level 3 load misses
PAPI_L3_STM	0x8000000f	No	No	Level 3 store misses
PAPI_BRU_IDL	0x80000010	No	No	Cycles branch units are idle
PAPI_FXU_IDL	0x80000011	No	No	Cycles integer units are idle
PAPI_FPU_IDL	0x80000012	No	No	Cycles floating point units are idle
PAPI_LSU_IDL	0x80000013	No	No	Cycles load/store units are idle
PAPI_TLB_DM	0x80000014	Yes	No	Data translation lookaside buffer misses
PAPI_TLB_IM	0x80000015	Yes	No	Instruction translation lookaside buffer misses
PAPI_TLB_TL	0x80000016	Yes	Yes	Total translation lookaside buffer misses
PAPI_L1_LDM	0x80000017	Yes	No	Level 1 load misses
PAPI_L1_STM	0x80000018	Yes	No	Level 1 store misses
PAPI_L2_LDM	0x80000019	Yes	No	Level 2 load misses
PAPI_L2_STM	0x8000001a	Yes	No	Level 2 store misses
PAPI_BTAC_M	0x8000001b	No	No	Branch target address cache misses
PAPI_PRF_DM	0x8000001c	No	No	Data prefetch cache misses

PAPI_L3_DCH	0x8000001d	No	No	Level 3 data cache hits
PAPI_TLB_SD	0x8000001e	No	No	Translation lookaside buffer shutdowns
PAPI_CSR_FAL	0x8000001f	No	No	Failed store conditional instructions
PAPI_CSR_SUC	0x80000020	No	No	Successful store conditional instructions
PAPI_CSR_TOT	0x80000021	No	No	Total store conditional instructions
PAPI_MEM_SCY	0x80000022	No	No	Cycles Stalled Waiting for memory accesses
PAPI_MEM_RCY	0x80000023	No	No	Cycles Stalled Waiting for memory Reads
PAPI_MEM_WCY	0x80000024	No	No	Cycles Stalled Waiting for memory writes
PAPI_STL_ICY	0x80000025	No	No	Cycles with no instruction issue
PAPI_FUL_ICY	0x80000026	No	No	Cycles with maximum instruction issue
PAPI_STL_CCY	0x80000027	No	No	Cycles with no instructions completed
PAPI_FUL_CCY	0x80000028	No	No	Cycles with maximum instructions completed
PAPI_HW_INT	0x80000029	No	No	Hardware interrupts
PAPI_BR_UCN	0x8000002a	Yes	No	Unconditional branch instructions
PAPI_BR_CN	0x8000002b	Yes	No	Conditional branch instructions
PAPI_BR_TKN	0x8000002c	Yes	No	Conditional branch instructions taken
PAPI_BR_NTK	0x8000002d	Yes	Yes	Conditional branch instructions not taken
PAPI_BR_MSP	0x8000002e	Yes	No	Conditional branch instructions mispredicted
PAPI_BR_PRC	0x8000002f	Yes	Yes	Conditional branch instructions correctly predicted
PAPI_FMA_INS	0x80000030	No	No	FMA instructions completed
PAPI_TOT_IIS	0x80000031	Yes	No	Instructions issued
PAPI_TOT_INS	0x80000032	Yes	No	Instructions completed
PAPI_INT_INS	0x80000033	No	No	Integer instructions
PAPI_FP_INS	0x80000034	Yes	No	Floating point instructions
PAPI_LD_INS	0x80000035	Yes	No	Load instructions
PAPI_SR_INS	0x80000036	Yes	No	Store instructions
PAPI_BR_INS	0x80000037	Yes	No	Branch instructions
PAPI_VEC_INS	0x80000038	No	No	Vector/SIMD instructions (could include integer)
PAPI_RES_STL	0x80000039	Yes	No	Cycles stalled on any resource
PAPI_FP_STAL	0x8000003a	No	No	Cycles the FP unit(s) are stalled
PAPI_TOT_CYC	0x8000003b	Yes	No	Total cycles
PAPI_LST_INS	0x8000003c	Yes	Yes	Load/store instructions completed
PAPI_SYC_INS	0x8000003d	No	No	Synchronization instructions completed

PAPI_L1_DCH	0x8000003e	Yes	Yes	Level 1 data cache hits
PAPI_L2_DCH	0x8000003f	Yes	Yes	Level 2 data cache hits
PAPI_L1_DCA	0x80000040	Yes	No	Level 1 data cache accesses
PAPI_L2_DCA	0x80000041	Yes	No	Level 2 data cache accesses
PAPI_L3_DCA	0x80000042	Yes	Yes	Level 3 data cache accesses
PAPI_L1_DCR	0x80000043	Yes	No	Level 1 data cache reads
PAPI_L2_DCR	0x80000044	Yes	No	Level 2 data cache reads
PAPI_L3_DCR	0x80000045	Yes	No	Level 3 data cache reads
PAPI_L1_DCW	0x80000046	Yes	No	Level 1 data cache writes
PAPI_L2_DCW	0x80000047	Yes	No	Level 2 data cache writes
PAPI_L3_DCW	0x80000048	Yes	No	Level 3 data cache writes
PAPI_L1_ICH	0x80000049	Yes	No	Level 1 instruction cache hits
PAPI_L2_ICH	0x8000004a	Yes	No	Level 2 instruction cache hits
PAPI_L3_ICH	0x8000004b	No	No	Level 3 instruction cache hits
PAPI_L1_ICA	0x8000004c	Yes	No	Level 1 instruction cache accesses
PAPI_L2_ICA	0x8000004d	Yes	No	Level 2 instruction cache accesses
PAPI_L3_ICA	0x8000004e	Yes	No	Level 3 instruction cache accesses
PAPI_L1_ICR	0x8000004f	Yes	No	Level 1 instruction cache reads
PAPI_L2_ICR	0x80000050	Yes	No	Level 2 instruction cache reads
PAPI_L3_ICR	0x80000051	Yes	No	Level 3 instruction cache reads
PAPI_L1_ICW	0x80000052	No	No	Level 1 instruction cache writes
PAPI_L2_ICW	0x80000053	No	No	Level 2 instruction cache writes
PAPI_L3_ICW	0x80000054	No	No	Level 3 instruction cache writes
PAPI_L1_TCH	0x80000055	No	No	Level 1 total cache hits
PAPI_L2_TCH	0x80000056	Yes	Yes	Level 2 total cache hits
PAPI_L3_TCH	0x80000057	No	No	Level 3 total cache hits
PAPI_L1_TCA	0x80000058	Yes	Yes	Level 1 total cache accesses
PAPI_L2_TCA	0x80000059	Yes	No	Level 2 total cache accesses
PAPI_L3_TCA	0x8000005a	Yes	No	Level 3 total cache accesses
PAPI_L1_TCR	0x8000005b	Yes	Yes	Level 1 total cache reads
PAPI_L2_TCR	0x8000005c	Yes	Yes	Level 2 total cache reads

PAPI_L3_TCR	0x80000005d	Yes	Yes	Level 3 total cache reads
PAPI_L1_TCW	0x80000005e	No	No	Level 1 total cache writes
PAPI_L2_TCW	0x80000005f	Yes	No	Level 2 total cache writes
PAPI_L3_TCW	0x800000060	Yes	No	Level 3 total cache writes
PAPI_FML_INS	0x800000061	No	No	Floating point multiply instructions
PAPI_FAD_INS	0x800000062	No	No	Floating point add instructions
PAPI_FDV_INS	0x800000063	No	No	Floating point divide instructions
PAPI_FSQ_INS	0x800000064	No	No	Floating point square root instructions
PAPI_FNV_INS	0x800000065	No	No	Floating point inverse instructions
PAPI_FP_OPS	0x800000066	Yes	Yes	Floating point operations
PAPI_SP_OPS	0x800000067	Yes	Yes	Floating point operations;
				optimized to count scaled single precision vector operations
PAPI_DP_OPS	0x800000068	Yes	Yes	Floating point operations;
				optimized to count scaled double precision vector operations
PAPI_VEC_SP	0x800000069	Yes	No	Single precision vector/SIMD instructions
PAPI_VEC_DP	0x80000006a	Yes	No	Double precision vector/SIMD instructions
PAPI_REF_CYC	0x80000006b	Yes	No	Reference clock cycles

 Of 108 possible events, 64 are available, of which 17 are derived.

papi_event_chooser Command

Not all events can be simultaneously monitored (at least not without multiplexing):

```
[rush:~]$ papi_event_chooser PRESET PAPI_FP_OPS
Event Chooser: Available events which can be added with given events.
-----
PAPI Version           : 5.1.1.0
Vendor string and code : GenuineIntel (1)
Model string and code  : Intel(R) Xeon(R) CPU E7- 4830  @ 2.13GHz (47)
CPU Revision           : 2.000000
CPUID Info             : Family: 6  Model: 47  Stepping: 2
CPU Max Megahertz      : 2127
CPU Min Megahertz      : 2127
Hdw Threads per core   : 1
Cores per Socket       : 8
NUMA Nodes             : 4
CPUs per Node          : 8
Total CPUs             : 32
Running in a VM         : no
Number Hardware Counters : 7
Max Multiplex Counters  : 64
-----

  Name      Code      Deriv Description (Note)
PAPI_L1_DCM 0x80000000 No   Level 1 data cache misses
PAPI_L1_ICM 0x80000001 No   Level 1 instruction cache misses
```

PAPI_L2_DCM	0x80000002	Yes	Level 2 data cache misses
PAPI_L2_ICM	0x80000003	No	Level 2 instruction cache misses
PAPI_L1_TCM	0x80000006	Yes	Level 1 cache misses
PAPI_L2_TCM	0x80000007	No	Level 2 cache misses
PAPI_L3_TCM	0x80000008	No	Level 3 cache misses
PAPI_L3_LDM	0x8000000e	No	Level 3 load misses
PAPI_TLB_DM	0x80000014	No	Data translation lookaside buffer misses
PAPI_TLB_IM	0x80000015	No	Instruction translation lookaside buffer misses
PAPI_TLB_TL	0x80000016	Yes	Total translation lookaside buffer misses
PAPI_L1_LDM	0x80000017	No	Level 1 load misses
PAPI_L1_STM	0x80000018	No	Level 1 store misses
PAPI_L2_LDM	0x80000019	No	Level 2 load misses
PAPI_L2_STM	0x8000001a	No	Level 2 store misses
PAPI_BR_UCN	0x8000002a	No	Unconditional branch instructions
PAPI_BR_CN	0x8000002b	No	Conditional branch instructions
PAPI_BR_TKN	0x8000002c	No	Conditional branch instructions taken
PAPI_BR_NTK	0x8000002d	Yes	Conditional branch instructions not taken
PAPI_BR_MSP	0x8000002e	No	Conditional branch instructions mispredicted
PAPI_BR_PRC	0x8000002f	Yes	Conditional branch instructions correctly predicted
PAPI_TOT_IIS	0x80000031	No	Instructions issued
PAPI_TOT_INS	0x80000032	No	Instructions completed
PAPI_FP_INS	0x80000034	No	Floating point instructions
PAPI_LD_INS	0x80000035	No	Load instructions
PAPI_SR_INS	0x80000036	No	Store instructions
PAPI_BR_INS	0x80000037	No	Branch instructions
PAPI_RES_STL	0x80000039	No	Cycles stalled on any resource
PAPI_TOT_CYC	0x8000003b	No	Total cycles

PAPI_LST_INS	0x8000003c	Yes	Load/store instructions completed
PAPI_L2_DCH	0x8000003f	Yes	Level 2 data cache hits
PAPI_L2_DCA	0x80000041	No	Level 2 data cache accesses
PAPI_L3_DCA	0x80000042	Yes	Level 3 data cache accesses
PAPI_L2_DCR	0x80000044	No	Level 2 data cache reads
PAPI_L3_DCR	0x80000045	No	Level 3 data cache reads
PAPI_L2_DCW	0x80000047	No	Level 2 data cache writes
PAPI_L3_DCW	0x80000048	No	Level 3 data cache writes
PAPI_L1_ICH	0x80000049	No	Level 1 instruction cache hits
PAPI_L2_ICH	0x8000004a	No	Level 2 instruction cache hits
PAPI_L1_ICA	0x8000004c	No	Level 1 instruction cache accesses
PAPI_L2_ICA	0x8000004d	No	Level 2 instruction cache accesses
PAPI_L3_ICA	0x8000004e	No	Level 3 instruction cache accesses
PAPI_L1_ICR	0x8000004f	No	Level 1 instruction cache reads
PAPI_L2_ICR	0x80000050	No	Level 2 instruction cache reads
PAPI_L3_ICR	0x80000051	No	Level 3 instruction cache reads
PAPI_L2_TCH	0x80000056	Yes	Level 2 total cache hits
PAPI_L2_TCA	0x80000059	No	Level 2 total cache accesses
PAPI_L3_TCA	0x8000005a	No	Level 3 total cache accesses
PAPI_L2_TCR	0x8000005c	Yes	Level 2 total cache reads
PAPI_L3_TCR	0x8000005d	Yes	Level 3 total cache reads
PAPI_L2_TCW	0x8000005f	No	Level 2 total cache writes
PAPI_L3_TCW	0x80000060	No	Level 3 total cache writes
PAPI_SP_OPS	0x80000067	Yes	Floating point operations; optimized to count scaled single precision
PAPI_DP_OPS	0x80000068	Yes	Floating point operations; optimized to count scaled double precision
PAPI_VEC_SP	0x80000069	No	Single precision vector/SIMD instructions
PAPI_VEC_DP	0x8000006a	No	Double precision vector/SIMD instructions
PAPI_REF_CYC	0x8000006b	No	Reference clock cycles

Total events reported: 57

event_choser.c PASSED

PAPI Examples

In this section we will work through a few simple examples of using the PAPI API, mostly focused on using the high-level API. And we will steer clear of native events, and leave those to tool developers.

Accessing Counters Through PAPI

Include files for constants and routine interfaces:

C: `papi.h`
F77: `f77papi.h`
F90: `f90papi.h`

PAPI Naming Scheme

The `C` interfaces:

PAPI C interface

(return type) PAPI_function_name(arg1, arg2, ...)

and `Fortran` interfaces

PAPI Fortran interfaces

PAPIF_function_name(arg1, arg2, ..., *check*)

note that the *check* parameter is the same type and value as the `C` return value.

Relation Between C and Fortran Types in PAPI

The following table shows the relation between the C and Fortran types used in PAPI:

Pseudo-type	Fortran type	Description
C_INT	INTEGER	Default Integer type
C_FLOAT	REAL	Default Real type
C_LONG_LONG	INTEGER*8	Extended size integer
C_STRING	CHARACTER*(PAPI_MAX_STR_LEN)	Fortran string
C_INT FUNCTION	EXTERNAL INTEGER FUNCTION	Fortran function returning integer result

High-level API Example in C

Let's consider the following example code for using the high-level API in C.

```
#include "papi.h"
#include <stdio.h>
#define NUM_EVENTS 2
#define THRESHOLD 10000
#define ERROR_RETURN(retval) { fprintf(stderr, "Error %d %s:line %d: \n", \
    retval, __FILE__, __LINE__); exit(retval); }

void computation_mult() { /* stupid codes to be monitored */
    double tmp=1.0;
    int i=1;
    for( i = 1; i < THRESHOLD; i++ ) {
        tmp = tmp*i;
    }
}

void computation_add() { /* stupid codes to be monitored */
    int tmp = 0;
    int i=0;

    for( i = 0; i < THRESHOLD; i++ ) {
        tmp = tmp + i;
    }
}
```

```
int main()
{
    /*Declaring and initializing the event set with the presets*/
    int Events[2] = {PAPI_TOT_INS, PAPI_TOT_CYC};
    /*The length of the events array should be no longer than the
       value returned by PAPI_num_counters.*/

    /*declaring place holder for no of hardware counters */
    int num_hwcnts = 0;
    int retval;
    char errstring[PAPI_MAX_STR_LEN];
    /*This is going to store our list of results*/
    long_long values[NUM_EVENTS];

    /*****
     * This part initializes the library and compares the version number of the
     * header file, to the version of the library, if these don't match then it
     * is likely that PAPI won't work correctly.If there is an error, retval
     * keeps track of the version number.
     *****/
    /*****
     * if((retval = PAPI_library_init(PAPI_VER_CURRENT)) != PAPI_VER_CURRENT ) {
     *     fprintf(stderr, "Error: %d %s\n",retval, errstring);
     *     exit(1);
     * }
     *****/
}
```

```
/* *****  
 * PAPI_num_counters returns the number of hardware counters the platform *  
 * has or a negative number if there is an error *  
 * ***** */  
if ((num_hwcntrs = PAPI_num_counters()) < PAPI_OK) {  
    printf("There are no counters available. \n");  
    exit(1);  
}  
  
printf("There are %d counters in this system\n", num_hwcntrs);  
  
/* *****  
 * PAPI_start_counters initializes the PAPI library (if necessary) and *  
 * starts counting the events named in the events array. This function *  
 * implicitly stops and initializes any counters running as a result of *  
 * a previous call to PAPI_start_counters. *  
 * ***** */  
  
if ( (retval = PAPI_start_counters(Events, NUM_EVENTS)) != PAPI_OK)  
    ERROR_RETURN(retval);  
  
printf("\nCounter Started: \n");  
  
/* Your code goes here*/  
computation_add();
```

```
/* *****  
 * PAPI_read_counters reads the counter values into values array *  
 * ***** */  
if ( (retval=PAPI_read_counters(values, NUM_EVENTS)) != PAPI_OK)  
    ERROR_RETURN(retval);  
printf("Read successfully\n");  
  
printf("The total instructions executed for addition are %lld \n",values[0]);  
printf("The total cycles used are %lld \n", values[1] );  
  
printf("\nNow we try to use PAPI_accum to accumulate values\n");  
  
/* Do some computation here */  
computation_add();  
/* *****  
 * What PAPI_accum_counters does is it adds the running counter values *  
 * to what is in the values array. The hardware counters are reset and *  
 * left running after the call. *  
 * ***** */  
if ( (retval=PAPI_accum_counters(values, NUM_EVENTS)) != PAPI_OK)  
    ERROR_RETURN(retval);  
  
printf("We did an additional %d times addition!\n", THRESHOLD);  
printf("The total instructions executed for addition are %lld \n",  
        values[0] );  
printf("The total cycles used are %lld \n", values[1] );
```

```
/* *****  
 * Stop counting events (this reads the counters as well as stops them) *  
***** */  
  
printf("\nNow we try to do some multiplications\n");  
computation_mult();  
  
/* ***** PAPI_stop_counters ***** */  
if ((retval=PAPI_stop_counters(values, NUM_EVENTS)) != PAPI_OK)  
    ERROR_RETURN(retval);  
  
printf("The total instruction executed for multiplication are %lld \n",  
       values[0] );  
printf("The total cycles used are %lld \n", values[1] );  
exit(0);  
}
```


Running on CCR:

```
[rush:/ifs/user/jonesm/d_papi]$ icc -I$PAPI/include -o highlev highlev.c -L$PAPI/lib -lpapi
[rush:/ifs/user/jonesm/d_papi]$ module list
Currently Loaded Modulefiles:
  1) null                2) modules              3) use.own              4) intel-mpi/4.1.1
  5) papi/v5.1.1         6) intel/13.1
[rush:/ifs/user/jonesm/d_papi]$ ./highlev
There are 7 counters in this system

Counter Started:
Read successfully
The total instructions executed for addition are 4406
The total cycles used are 8482

Now we try to use PAPI_accum to accumulate values
We did an additional 10000 times addition!
The total instructions executed for addition are 10056
The total cycles used are 16570

Now we try to do some multiplications
The total instruction executed for multiplication are 6041
The total cycles used are 7328
```

PAPI Initialization

The preceding example used `PAPI_library_init` to initialize PAPI, which is also used for the low-level API, but you can also use the `PAPI_num_counters`, `PAPI_start_counters`, or one of the “rate” calls, `PAPI_flips`, `PAPI_flops`, or `PAPI_ipc`.

Events are counted, as we saw in the example, using `PAPI_accum_counters`, `PAPI_read_counters`, and `PAPI_stop_counters`.

Let's look at an even simpler example just using one of the rate counters.

High-level Example in F90

For something a little different we can look at our old friend, matrix multiplication, this time in Fortran:

```
! A simple example for the use of PAPI, the number of flops you should
! get is about INDEX^3 on machines that consider add and multiply one flop
! such as SGI, and 2*(INDEX^3) that don't consider it 1 flop such as INTEL
! -Kevin London
program flops
  implicit none
  include "f90papi.h"

  integer,parameter :: i8=SELECTED_INT_KIND(16) ! integer*8
  integer,parameter :: index=1000
  real :: matrixa(index,index),matrixb(index,index),mres(index,index)
  real :: proc_time, mflops, real_time
  integer(kind=i8) :: flpins
  integer :: i,j,k, retval
```

```
retval = PAPI_VER_CURRENT
CALL PAPIf_library_init(retval)
if ( retval.NE.PAPI_VER_CURRENT) then
  print*, 'Failure in PAPI_library_init: ', retval
end if

CALL PAPIf_query_event(PAPI_FP_OPS, retval)
if (retval .NE. PAPI_OK) then
  print*, 'Sorry, no PAPI_FP_OPS event: ', PAPI_ENOEVNT
end if

! Initialize the Matrix arrays
do i=1,index
  do j=1,index
    matrixa(i,j) = i+j
    matrixb(i,j) = j-i
    mres(i,j) = 0.0
  end do
end do

! Setup PAPI library and begin collecting data from the counters
call PAPIf_flops( real_time, proc_time, flpins, mflops, retval )
if ( retval.NE.PAPI_OK) then
  print*, 'Failure on PAPIf_flops: ', retval
end if
```

```
! Matrix-Matrix Multiply
do i=1,index
  do j=1,index
    do k=1,index
      mres(i,j) = mres(i,j) + matrixa(i,k)*matrixb(k,j)
    end do
  end do
end do

! Collect the data into the Variables passed in
call PAPIf_flops( real_time, proc_time, flpins, mflops, retval)
if ( retval.NE.PAPI_OK ) then
  print*, 'Failure on PAPIf_flops: ', retval
end if
print *, 'Real_time: ', real_time
print *, ' Proc_time: ', proc_time
print *, ' Total flpins: ', flpins
print *, ' MFLOPS: ', mflops
end program flops
```

Compile and run on E5645 U2:

```
[rush:~]$ ifort -I$PAPI/include -o flops flops.f90 -L$PAPI/lib -lpapi
[rush:/ifs/user/jonesm/d_papi]$ ./flops
Real_time:    0.4914970
Proc_time:    0.4900359
Total flpins:                503306816
MFLOPS:       1027.082
```

Low-level API

The low-level API is primarily intended for experienced application programmers and tool developers. It manages hardware events in user-defined groups called “event sets,” and can use both preset and native events. The low-level API can also interrogate the hardware and determine memory sizes of the executable itself.

The low-level API can also be used for **multiplexing**, in which more (virtual) counters can be used than the underlying hardware supports by timesharing the available (physical) hardware counters.

PAPI Low-level Example

A simple example using the low-level API:

```
#include <papi.h>
#include <stdio.h>

#define NUM_FLOPS 10000

main() {
    int retval, EventSet=PAPI_NULL;
    long_long values[1];

    /* Initialize the PAPI library */
    retval = PAPI_library_init(PAPI_VER_CURRENT);
    if (retval != PAPI_VER_CURRENT) {
        fprintf(stderr, "PAPI library init error!\n");
        exit(1);
    }

    /* Create the Event Set */
    if (PAPI_create_eventset(&EventSet) != PAPI_OK) handle_error(1);

    /* Add Total Instructions Executed to our Event Set */
    if (PAPI_add_event(EventSet, PAPI_TOT_INS) != PAPI_OK) handle_error(1);
```



```
/* Start counting events in the Event Set */
if (PAPI_start(EventSet) != PAPI_OK) handle_error(1);

/* Defined in tests/do_loops.c in the PAPI source distribution */
do_flops(NUM_FLOPS);

/* Read the counting events in the Event Set */
if (PAPI_read(EventSet, values) != PAPI_OK) handle_error(1);

printf("After reading the counters: %lld\n", values[0]);

/* Reset the counting events in the Event Set */
if (PAPI_reset(EventSet) != PAPI_OK) handle_error(1);

do_flops(NUM_FLOPS);

/* Add the counters in the Event Set */
if (PAPI_accum(EventSet, values) != PAPI_OK) handle_error(1);
printf("After adding the counters: %lld\n", values[0]);

do_flops(NUM_FLOPS);

/* Stop the counting of events in the Event Set */
if (PAPI_stop(EventSet, values) != PAPI_OK) handle_error(1);

printf("After stopping the counters: %lld\n", values[0]);
}
```

PAPI in Parallel

threads `PAPI_thread_init` enables PAPI's thread support, and should be called immediately after `PAPI_library_init`.

MPI codes are treated very simply - each process has its own address space, and potentially its own hardware counters.

High-level Tools

There are a bunch of open-source high-level tools that build on some of the simple approaches that we have been talking about. General characteristics found in most (not necessarily all):

- Ability to generate and view MPI trace files, leveraging MPI's built-in profiling interface,
- Ability to do statistical profiling (à la `gprof`) and code viewing for identifying hotspots,
- Ability to access performance counters, leveraging PAPI

Tool Examples

A list of such high-level tool examples (not exhaustive):

- **TAU**, Tuning and Analysis Utility,

<http://www.cs.uoregon.edu/Research/tau/home.php>

- **Open|SpeedShop**, funded by U.S. DOE

<http://www.openspeedshop.org>

- **IPM**, Integrated Performance Management

<http://ipm-hpc.sourceforge.net>

Example: IPM

IPM is relatively simple to install and use, so we can easily walk through our favorite example. Note that IPM does:

- MPI
- PAPI
- I/O profiling
- Memory
- Timings: wall, user, and system

Run and Gather

```
1  #!/bin/bash
2  #SBATCH --nodes=1
3  #SBATCH --ntasks-per-node=16
4  #SBATCH --constraint=CPU-E5-2660
5  #SBATCH --partition=debug
6  #SBATCH --time=00:15:00
7  #SBATCH --mail-type=END
8  #SBATCH --mail-user=jonesm@buffalo.edu
9  #SBATCH --output=slurmIPM.out
10 #SBATCH --job-name=ipm-test
11 module load intel
12 module load intel-mpi
13 module list
14 export I_MPI_DEBUG=4
15 # Use LD_PRELOAD trick to load ipm wrappers at runtime
16 export LD_PRELOAD=/projects/jonesm/ipm/src/ipm/lib/libipm.so
17 export I_MPI_PMI_LIBRARY=/usr/lib64/libpmi.so
18 srun ./laplace_mpi<<EOF
19 2000
20 EOF
```

... and the output is a big xml file plus some useful output to standard output:

```

1 [rush:~/d_laplace/d_ipm]$ file jonesm.1318862245.001449.0
2 jonesm.1318862245.001449.0: XML document text
3 [rush:~/d_laplace/d_ipm]$ less subMPIP.out
4 ...
5 ##IPMv0.983#####
6 #
7 # command : ./laplace_mpi (completed)
8 # host      : d16n03/x86_64_Linux          mpi_tasks : 16 on 2 nodes
9 # start     : 10/17/11/10:37:25            wallclock : 116.170005 sec
10 # stop      : 10/17/11/10:39:21            %comm      : 13.94
11 # gbytes    : 2.24606e+00 total            gflop/sec   : 5.02520e+00 total
12 #
13 #####
14 # region : *          [ntasks] =      16
15 #

```

```

1  #                [total]                <avg>                min                max
2  # entries                16                1                1                1
3  # wallclock                1853.71                115.857                115.816                116.17
4  # user                1853.09                115.818                115.707                115.936
5  # system                2.18066                0.136291                0.071989                0.198969
6  # mpi                259.152                16.197                11.3859                19.1157
7  # %comm                13.9425                9.82914                16.5048
8  # gflop/sec                5.0252                0.314075                0.311741                0.319497
9  # gbytes                2.24606                0.140379                0.138138                0.170021
10 #
11 # PAPI_FP_OPS                5.83778e+11                3.64861e+10                3.62149e+10                3.7116e+10
12 # PAPI_FP_INS                5.8276e+11                3.64225e+10                3.62144e+10                3.69079e+10
13 # PAPI_DP_OPS                5.82764e+11                3.64228e+10                3.62144e+10                3.69079e+10
14 # PAPI_VEC_DP                4.00803e+06                250501                0                4.00803e+06
15 #
16 #                [time]                [calls]                <%mpi>                <%wall>
17 # MPI_Allreduce                243.838                381520                94.09                13.15
18 # MPI_Sendrecv                14.9598                763040                5.77                0.81
19 # MPI_Send                0.339084                15                0.13                0.02
20 # MPI_Recv                0.0143731                15                0.01                0.00
21 # MPI_Bcast                0.00124932                16                0.00                0.00
22 # MPI_Comm_rank                1.58967e-05                16                0.00                0.00
23 # MPI_Comm_size                8.01496e-06                16                0.00                0.00
24 #####

```

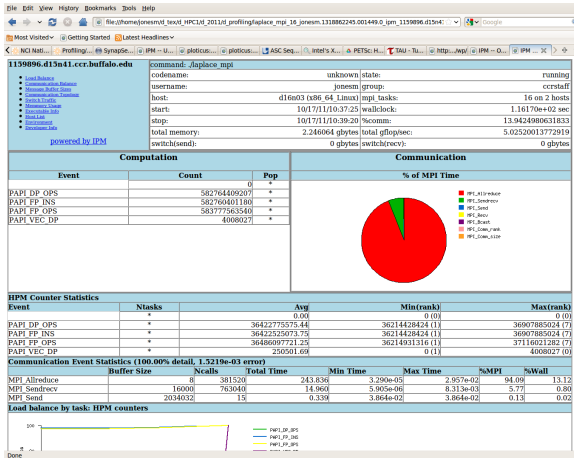

Script to Generate HTML from XML Results

```
1  #!/bin/bash
2
3  if [ $# -ne 1 ]; then
4      echo "Usage: $0 xml_filename"
5      exit
6  fi
7  XMLFILE=$1
8
9  export IPM_KEYFILE=/projects/jonesm/ipm/src/ipm/ipm_key
10 export PATH=${PATH}:/projects/jonesm/ipm/src/ipm/bin
11 /projects/jonesm/ipm/src/ipm/bin/ipm_parse -html $XMLFILE
```

```
[u2:~/d_laplace/d_ipm]$ ./genhtml.sh jonesm.1318862245.001449.0
# data_acquire      = 0 sec
# data_workup       = 0 sec
# mpi_pie           = 1 sec
# task_data         = 0 sec
# load_bal          = 0 sec
# time_stack        = 0 sec
# mpi_stack         = 1 sec
# mpi_buff          = 0 sec
# switch+mem        = 0 sec
# topo_tables       = 0 sec
# topo_data         = 0 sec
# topo_time         = 0 sec
# html_all          = 2 sec
# html_regions      = 0 sec
# html_nonregion    = 1 sec
[rush:~/d_laplace/d_ipm]$ ls -l \
laplace_mpi_16_jonesm.1318862245.001449.0_ipm_1159896.d15n41.ccr.buffalo.edu/
total 346
-rw-r--r-- 1 jonesm ccrstaff  994 Oct 17 16:07 dev.html
-rw-r--r-- 1 jonesm ccrstaff  104 Oct 17 16:07 env.html
-rw-r--r-- 1 jonesm ccrstaff  347 Oct 17 16:07 exec.html
-rw-r--r-- 1 jonesm ccrstaff  451 Oct 17 16:07 hostlist.html
drwxr-xr-x 2 jonesm ccrstaff  930 Oct 17 16:07 img
-rw-r--r-- 1 jonesm ccrstaff 10550 Oct 17 16:07 index.html
-rw-r--r-- 1 jonesm ccrstaff   387 Oct 17 16:07 map_adjacency.txt
-rw-r--r-- 1 jonesm ccrstaff  8961 Oct 17 16:07 map_calls.txt
-rw-r--r-- 1 jonesm ccrstaff  1452 Oct 17 16:07 map_data.txt
drwxr-xr-x 2 jonesm ccrstaff   803 Oct 17 16:07 pl
-rw-r--r-- 1 jonesm ccrstaff  2620 Oct 17 16:07 task_data
[rush:~/d_laplace/d_ipm]$ tar czf my-ipm-files.tgz \
laplace_mpi_16_jonesm.1318862245.001449.0_ipm_1159896.d15n41.ccr.buffalo.edu/
[rush:~/d_laplace/d_ipm]$ ls -l my-ipm-files.tgz
-rw-r--r-- 1 jonesm ccrstaff 71509 Oct 17 16:48 my-ipm-files.tgz
```

Visualize Results in Browser

Transfer compressed tar file to your local machine, unpack, and browse the results:



Summary

Summary of high-level tools

- IPM is pretty easy to use, provides some good functionality
- TAU and Open|SpeedShop have steeper learning curves, much more functionality