

# ASSIGNMENT 4 SOLUTIONS

## HPC1 Fall 2013

**Due Date:** *Tuesday, November 5*

(please submit your report electronically to the instructor, in one PDF file, as *hw4-yourUBitname.pdf*)

**Problem 1:** Consider the Laplace equation in two dimensions (we will consider Dirichlet boundary conditions in the unit interval),

$$\nabla^2 \phi = \left( \frac{\partial^2}{\partial x^2} + \frac{\partial^2}{\partial y^2} \right) \phi = 0. \quad (1)$$

Now take the case of a uniform square grid of  $M$  points in each direction, and using a simple two-point numerical derivative, it is not difficult to show that the optimum value of  $\phi$  at the point  $(i, j) = (x_i, y_j)$  is given by:

$$\phi_{ij} \simeq \frac{1}{4} (\phi_{i+1,j} + \phi_{i,j+1} + \phi_{i-1,j} + \phi_{i,j-1}) \quad (i, j) \in (1..M, 1..M). \quad (2)$$

The *relaxation method* consists of iterating this equation to obtain a new, improved solution from the previous iteration:

$$\phi_{ij}^{n+1} \simeq \frac{1}{4} (\phi_{i+1,j}^n + \phi_{i,j+1}^n + \phi_{i-1,j}^n + \phi_{i,j-1}^n) \quad (i, j) \in (1..M, 1..M) \quad (3)$$

A simple summary of the technique (otherwise known as Jacobi iteration):

1. Apply a square lattice with uniform spacing - label the points  $(i, j)$ .
  2. Apply the fixed boundary condition values.
  3. Make an initial guess for the interior points  $\phi_{i,j}^0$ .
  4. Iterate until convergence, using the “cross” scheme in Eq. 3 (note that better schemes are available - feel free to derive one, or look them up in standard references).
- a.** Write a serial code to perform the above solution for the problem on the unit square ( $0 \leq x, y \leq 1$ ) with boundary conditions

$$\begin{aligned} \phi(x, 0) &= \sin(\pi x), \\ \phi(x, 1) &= \sin(\pi x)e^{-\pi}, \\ \phi(0, y) &= \phi(1, y) = 0. \end{aligned}$$

You can verify that the analytic solution to this problem is  $\phi(x, y) = \sin(\pi x)e^{-\pi y}$ .

**Solution:**

Solution combined with part **c** with a code listing included at the end of these solutions.

- b. How does the number of Jacobi iterations required for convergence in your code depend on the grid size? Show a plot of this behavior. Take a reasonable threshold for convergence, say to where the L2-norm of the solution difference between iterations reaches  $10^{-5}$ .

*Solution:*




Figure 1: Number of iterations required as a function of the grid size for the simple Jacobi solver, using an  $L_2$  norm of  $10^{-4}$  and  $10^{-5}$ .

Plotted in Fig. 1 is the number of Jacobi iterations as a function of the grid size, for a tolerance of  $10^{-4}$  and  $10^{-5}$  (a stricter tolerance will require more iterations) of the  $L_2$ -norm between the solution iterates. Note the strong nonlinear growth (in fact proportional to the number of grid points squared) owing to the nature of relaxation methods - the solution actually propagates inward from the domain boundary, so as the mesh size grows, so must the number of iterations, roughly as the square of the number of mesh points. Note that the error, as measured from the analytic solution, is considerably larger than this iterative tolerance (to improve the error, reduce the tolerance).

You should verify, of course, that you are getting a “good” solution relative to the analytic one - see Fig. 2 below (a good solution will require an even stricter convergence criterion).




Figure 2: Color contour of solution using Jacobi solver, plotted using IDL.

- c. Parallelize your solver using OpenMP.

*Solution:*

See the program listing at the end of the solutions.

- d. Take a grid size that is significant enough to require some time (say 1000 or so, but feel free to innovate), and do a study in parallel scalability - how well does your parallel program scale with additional processors?

*Solution:*




Figure 3: OpenMP parallel performance of the simple Jacobi solver on the **UB/CCR** nodes (12-core systems).

The performance shown in Fig. 3 is not bad - a speedup of about 8 on all 12 cores. Note the variation shown by adjusting the affinity option (which in the Intel compiler is done by setting the `KMP_AFFINITY` environment variable) such that the OpenMP threads are placed on adjoining cores (**compact** setting), or filled in using a scatter approach (**scatter** setting) on the other processors (recall that these 12-core nodes have two 6-core processors, with a shared L3-cache on each processor). What we can see from the difference between the **compact** and **scatter** approaches is there is very significant contention for memory in this code (since our matrices are 1024x1024, each one takes up about 8MB of memory, so the two matrices together can not fit into the L3 cache). Using the **scatter** approach for thread placement gives a performance boost for intermediate values of `OMP_NUM_THREADS`, as the second processor's L3 cache comes into use, but once the node is full, a **compact** strategy will actually work best. What happens if you fit both matrices into the L3 cache? You can see from Fig. 4 the

results from running the same grid size on one of the Intel 32-core nodes, which have 24MB of L3 cache for each of the four (8 cores each) processors. If you increase the problem size

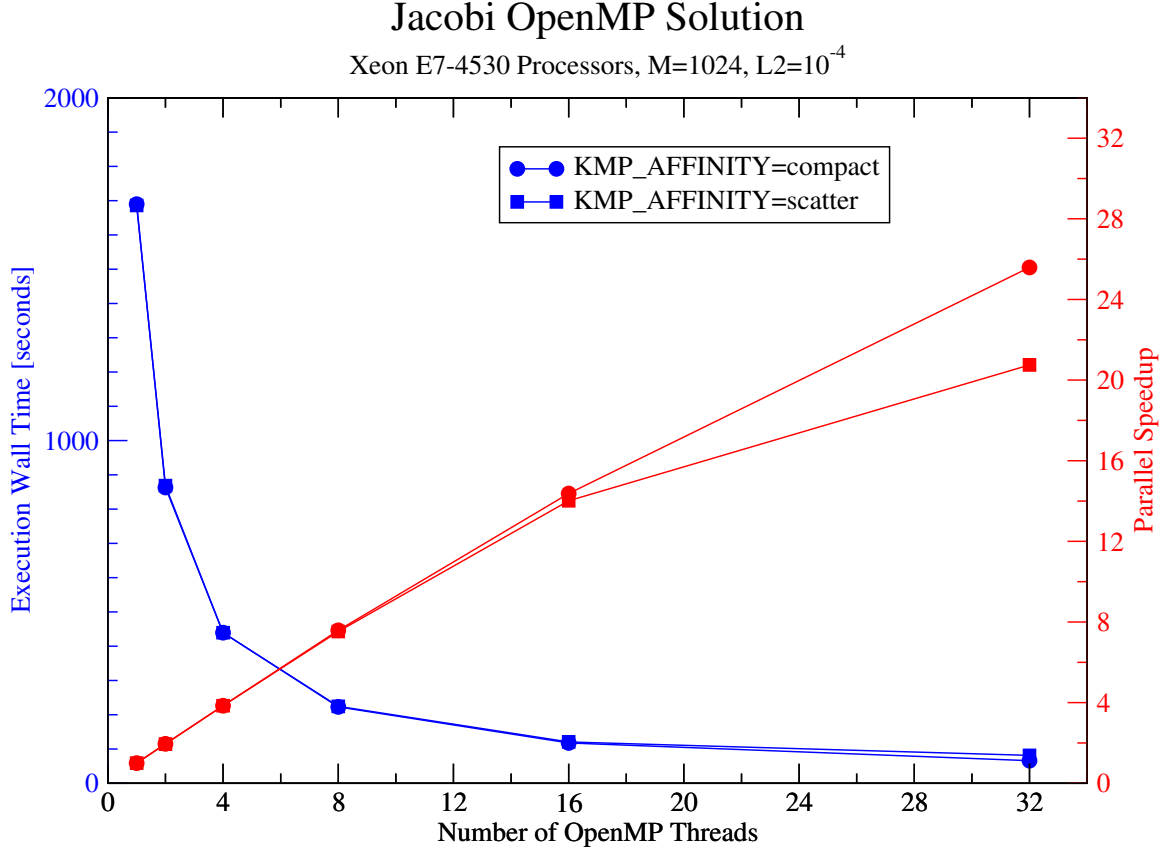


Figure 4: OpenMP parallel performance of the simple Jacobi solver on the **UB/CCR** nodes (32-core systems).

still further, however, the same behavior that we saw in Fig. 3 will manifest itself as the matrices grow too large to completely fit into the cache, and the cores contend for access to main memory.

```

1  MODULE serialMOD
2      implicit none
3      integer,parameter :: dp=selected_real_kind(2*PRECISION(1.0))
4      real(kind=dp),parameter :: PI=3.1415926535897932
5
6  CONTAINS
7      subroutine BC1(sol,m)
8          !
9          ! sol = current solution
10         ! m = 2d grid extent (square assumed 0:m+1)
11         !
12         ! BC: 0<= x,y <=1
13         ! y=0: sol = sin(\pi*x)
14         ! y=1: sol = exp(-\pi)*sin(\pi*x)
15         !
16         implicit none
17         integer :: m
18         real(kind=dp) :: sol(0:m+1,0:m+1)
19
20         integer :: i
21         real(kind=dp) :: x
22
23         ! BCs initialized to Dirichlet values, rest of points
24         ! set to zero
25         sol=0.0_dp
26         do i=0,m+1
27             x=dbl(i)/dbl(m+1)
28             sol(i,0)=sin(PI*x)
29             sol(i,m+1)=sin(PI*x)*exp(-PI)
30         end do
31         return
32     end subroutine BC1
33 end MODULE serialMOD
34
35 program laplace_omp
36     USE serialMOD
37     USE omp_lib
38     implicit none
39
40     integer,parameter :: MAXITER=15000000,outiter=10000
41     !real(kind=dp),parameter :: tolerance=1.0d-6
42     real(kind=dp) :: tolerance
43     integer :: i,j,m,iter,nthreads,ctick1,ctick2,ctickrate,ctickmax
44     real(kind=dp) :: t0,t1,t2,t3,fops,delta,tdelta,residual,exacterr,exactsol
45     real(kind=dp),allocatable :: sol(:,,:),oldsol(:,,:)
46     real(kind=dp) :: dsecnd
47     external dsecnd
48
49     print*,'Enter grid extent, m: (0:m+1) and tolerance'
50     read(*,*) m,tolerance
51     print*,'Using grid of size: ',m,', 0,m+1 used for BCs.'
52     print*,'Using Tolerance = ',tolerance
53
54     ALLOCATE(sol(0:m+1,0:m+1),oldsol(0:m+1,0:m+1))
55     call BC1(sol,m) ! Apply Dirichlet boundary conditions
56     oldsol=sol
57
58 #ifdef _OPENMP
59     t0=OMP_GET_WTIME()
60 #else
61     t0=dsecnd()
62 #endif
63     iter=0
64     residual=1.d6
65     do while ( (residual > tolerance).and.(iter <= MAXITER) )
66         iter = iter+1
67         ! update interior points based on "+" points
68         residual = 0.0_dp
69         !$OMP PARALLEL DEFAULT(SHARED) PRIVATE(i,j,delta)
70         !$OMP DO REDUCTION(+:residual)
71         do j=1,m
72             do i=1,m
73                 sol(i,j)=0.25_dp*(oldsol(i+1,j)+oldsol(i-1,j)+oldsol(i,j+1)+oldsol(i,j-1))
74                 delta = sol(i,j)-oldsol(i,j)
75                 residual = residual+delta*delta

```

```

76         end do
77     end do
78 !$OMP END DO
79 !$OMP DO
80     do j=1,m
81         do i=1,m
82             oldsol(i,j)=sol(i,j)
83         end do
84     end do
85 !$OMP END DO
86 !$OMP END PARALLEL
87     residual = sqrt(residual)
88     if(iter.eq.1.or.MOD(iter,outiter).eq.0) &
89         write(*, '(' iter = ",i6," residual = ",e12.4)') iter,residual
90 end do
91 fops = 8.0*m*m
92 fops = fops*iter
93 #ifdef _OPENMP
94     t1=OMP_GET_WTIME()
95 #else
96     t1=dsecnd()
97 #endif
98     write(*, '('Total Runtime[s], (under)estimated Mflop/s: ",f14.6,f12.2)') t1-t0, &
99         fops/(1.e6*(t1-t0))
100 !
101 ! Output results for plotting
102 !
103 exacterr = -1.0
104 do i=0,m+1
105     do j=0,m+1
106         exactsol = SIN( PI*DBLE(i)/DBLE(m+1) )*EXP( -PI*DBLE(j)/DBLE(m+1) )
107         exacterr = MAX( exacterr, ABS(exactsol-sol(i,j)) )
108 !         print*, 'i,j,err = ',i,j,exactsol-sol(i,j)
109     end do
110 end do
111 open(unit=10,file="laplace_omp.dat",status='unknown',form='unformatted')
112 print*, ' Max value in sol: ',MAXVAL(sol)
113 print*, ' Min value in sol: ',MINVAL(sol)
114 print*, ' Max error from analytic solution: ',exacterr
115 write(10) ((sngl(sol(i,j)),i=0,m+1),j=0,m+1)
116 close(10)
117 DEALLOCATE(sol,oldsol)
118 print*, 'Solution reached in ',iter,' iterations.'
119 end program laplace_omp

```